

LAPORAN TUGAS 2

1. RINGKASAN EKSEKUTIF

Tugas ini mengimplementasikan Distributed Synchronization System yang mensimulasikan skenario realworld dari distributed systems. Sistem ini mampu menangani multiple nodes (minimal 3) yang berkomunikasi dan mensinkronisasi data secara konsisten dengan mekanisme fault tolerance dan recovery.

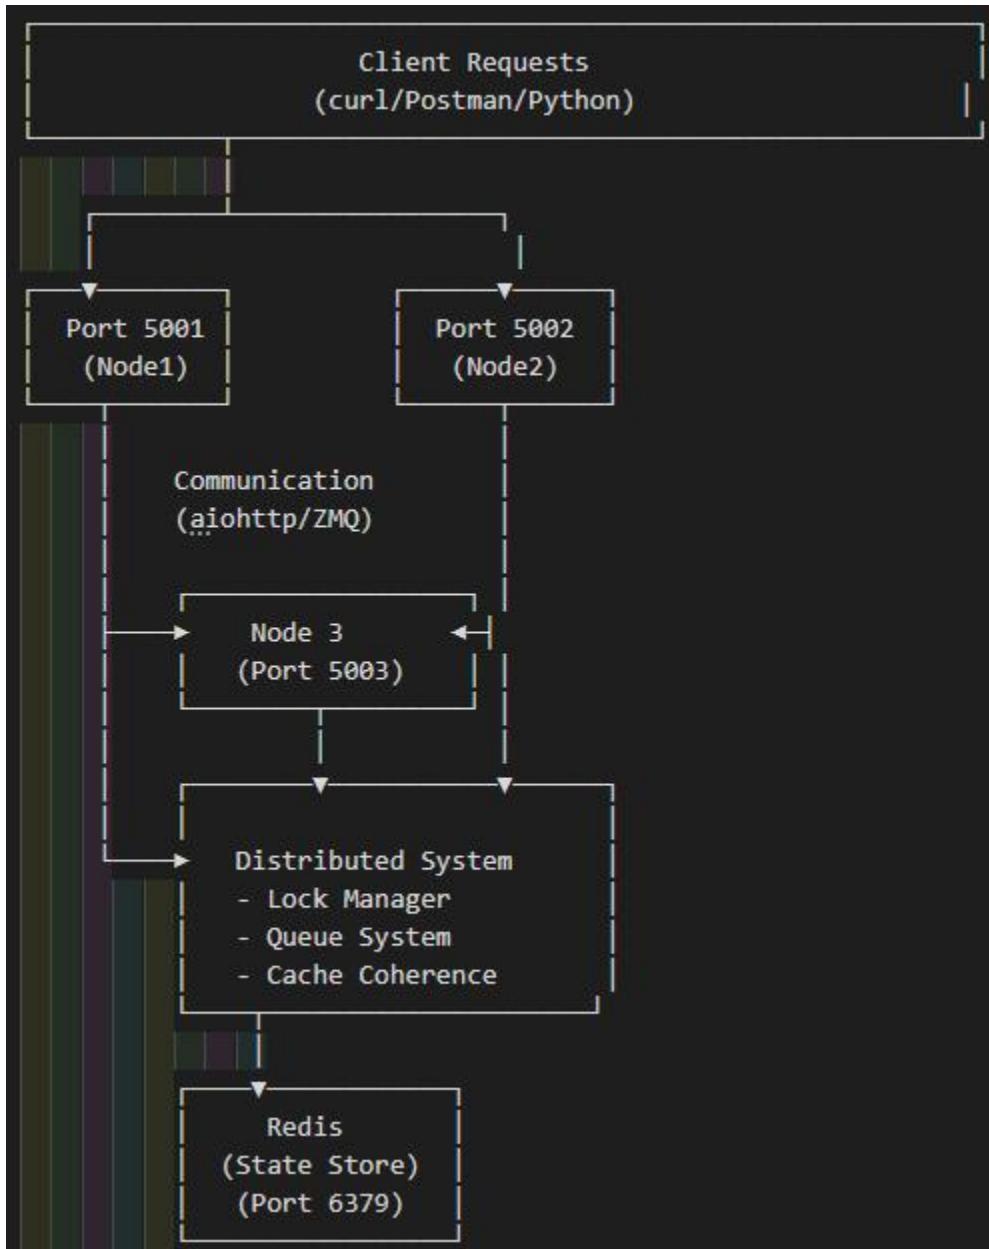
Komponen Utama:

1. Distributed Lock Manager Menggunakan Raft Consensus
2. Distributed Queue System Menggunakan Consistent Hashing
3. Distributed Cache Coherence Protokol MESI/MOSI/MOESI
4. Containerization Docker & Docker Compose

2. TECHNICAL DOCUMENTATION

2.1 Arsitektur Sistem Lengkap

Diagram Arsitektur Sistem



Komponen Sistem

Komponen	Deskripsi	Teknologi
Lock Manager	Mengelola distributed locks dengan Raft consensus	Python asyncio, aiohttp
Queue System	Mendistribusikan pesan dengan consistent hashing	Redis, Python

Cache Node	Mengelola cache lokal dengan protocol coherence	In-memory cache, Python
Communication Layer	Internode communication	aiohttp, ZeroMQ
State Store	Persistent storage untuk distributed state	Redis
Monitoring	Performance metrics dan logging	Python logging, Prometheus

2.2 Penjelasan Algoritma yang Digunakan

A. Raft Consensus Algorithm

Tujuan: Memastikan konsistensi data di antara multiple nodes dalam kondisi faulttolerant.

Komponen Utama:

- Term: Epoch/periode dalam Raft, dimulai dari 0 dan terus meningkat
- Log Entry: Unit basic dari replikasi, berisi command dan term
- State Machine: Aplikasi logic yang diterapkan pada log yang tercommit

Mekanisme Raft:

1. Election (Pemilihan Leader)

Setiap node dimulai sebagai FOLLOWER

Jika FOLLOWER tidak menerima heartbeat dalam election timeout, menjadi CANDIDATE
CANDIDATE mengirim RequestVote RPC ke semua nodes

Node yang menerima majority votes menjadi LEADER untuk term saat ini

2. Log Replication (Replikasi Log)

LEADER menerima client requests dan membuat log entries

LEADER mengirim AppendEntries RPC ke semua FOLLOWERS

FOLLOWERS menambahkan entries ke log mereka
Setelah majority replies, LEADER commits entries
LEADER menginformasikan FOLLOWERS tentang committed entries

3. Safety

Election Safety: Hanya satu LEADER per term

Log Matching: Jika dua log entries memiliki term dan index sama, maka semua entries sebelumnya juga identik

State Machine Safety: Jika state machine menerapkan log entry pada index tertentu, tidak ada state machine lain yang menerapkan entry berbeda pada index yang sama

Implementasi di Sistem:

```
'''python
```

class RaftNode:

```
def __init__(self, node_id, peers):
    self.node_id = node_id
    self.peers = peers
    self.state = 'FOLLOWER' FOLLOWER, CANDIDATE, LEADER
    self.term = 0
    self.voted_for = None
    self.log = []
    self.commit_index = 0
    self.last_applied = 0
...
...
```

B. Consistent Hashing (Distributed Queue)

Tujuan: Mendistribusikan messages ke nodes dengan minimal redistribution saat node join/leave.

Mekanisme:

1. Hash ring dengan virtual nodes
2. Message key dihash dan ditempatkan pada node terdekat
3. Jika node failure, hanya messages di node tersebut yang direbalance

Keuntungan:

Scalability: Minimal data movement saat scaling

Fault Tolerance: Data tetap accessible melalui node lain (replication)

Load Balancing: Messages didistribusikan merata

C. Cache Coherence Protocol (MESI)

State Cache Line:

Modified (M): Cache line telah dimodifikasi, berbeda dari main memory

Exclusive (E): Cache line clean, tidak ada copy di cache lain

Shared (S): Cache line clean, ada copy di cache lain

Invalid (I): Cache line invalid

Transisi State:

Current	Action	New State
I	Read Miss	S/E
I	Write Miss	M
E	Read	E
E	Write	M
S	Read	S
S	Write	Invalidate others, M
M	Read	Other S + Writeback
M	Write	Other Invalidate

Replacement Policy: LRU (Least Recently Used)

Cache line yang paling lama tidak diakses akan dievict terlebih dahulu

Maintain queue dari leastrecent ke mostrecent

2.3 API Documentation dengan OpenAPI/Swagger Spec

Base URL: `http://localhost:5001` (untuk node1), etc.

Lock Manager Endpoints

```
```yaml
```

```
/lock/acquire:
```

```
post:
```

```
summary: Acquire lock untuk resource tertentu
```

```
requestBody:
```

```
required: true
```

```
content:
```

```
application/json:
```

```
schema:
```

```
type: object
```

```
properties:
```

```
resource:
```

```
type: string
```

```
example: "database_connection"
```

```
mode:
```

```
type: string
```

```
enum: [shared, exclusive]
```

```
example: "exclusive"

timeout:
 type: integer
 description: Timeout dalam detik (opsional)
 example: 30

responses:
 '200':
 description: Lock berhasil diperoleh
 content:
 application/json:
 schema:
 type: object
 properties:
 status: string
 lock_id: string
 acquired_at: string
 '409':
 description: Lock conflict atau timeout
 '500':
 description: Server error

/lock/release:
 post:
 summary: Release lock yang telah diperoleh
 requestBody:
 required: true
 content:
 application/json:
```

```
schema:
 type: object
 properties:
 resource:
 type: string
 lock_id:
 type: string
responses:
 '200':
 description: Lock berhasil dirilis
 '404':
 description: Lock tidak ditemukan
 '500':
 description: Server error

/lock/status:
get:
 summary: Cek status lock untuk resource
 parameters:
 in: query
 name: resource
 required: true
 schema:
 type: string
responses:
 '200':
 description: Status lock
 content:
```

```
application/json:
 schema:
 type: object
 properties:
 resource: string
 locked: boolean
 holder: string
 mode: string
 ...
```

## Queue Endpoints

```
```yaml  
/queue/enqueue:  
  post:  
    summary: Masukkan pesan ke queue  
    requestBody:  
      required: true  
      content:  
        application/json:  
          schema:  
            type: object  
            properties:  
              message:  
                type: string  
                example: "Task data"  
              priority:  
                type: integer
```

```
    minimum: 1
    maximum: 10
  responses:
    '200':
      description: Pesan berhasil masuk queue
      content:
        application/json:
          schema:
            type: object
            properties:
              message_id: string
              enqueueued_at: string
    '500':
      description: Server error

/queue/dequeue:
  get:
    summary: Ambil pesan dari queue
  responses:
    '200':
      description: Pesan berhasil diambil
      content:
        application/json:
          schema:
            type: object
            properties:
              message_id: string
              message: string
```

```
    dequeued_at: string  
'204':  
    description: Queue kosong  
'500':  
    description: Server error
```

/queue/status:

```
get:  
summary: Cek status queue  
responses:  
'200':  
    description: Status queue  
    content:  
        application/json:  
        schema:  
            type: object  
            properties:  
                queue_length: integer  
                pending_messages: array  
...  
...
```

Cache Endpoints

```
'''yaml  
/cache/set:  
post:  
summary: Set nilai di cache  
requestBody:
```

```
required: true  
content:  
  application/json:  
    schema:  
      type: object  
      properties:  
        key:  
          type: string  
        value:  
          type: string  
        ttl:  
          type: integer  
          description: Time to live dalam detik  
responses:  
  '200':  
    description: Nilai berhasil diset  
  '500':  
    description: Server error
```

```
/cache/get:  
get:  
  summary: Ambil nilai dari cache  
parameters:  
  in: query  
  name: key  
  required: true  
schema:  
  type: string
```

```
responses:  
  '200':  
    description: Nilai ditemukan  
    content:  
      application/json:  
        schema:  
          type: object  
          properties:  
            key: string  
            value: string  
            cache_state: string  
  '404':  
    description: Key tidak ditemukan  
  '500':  
    description: Server error
```

```
/cache/invalidate:  
  post:  
    summary: Invalidate cache entry  
    requestBody:  
      required: true  
      content:  
        application/json:  
          schema:  
            type: object  
            properties:  
              key:  
                type: string
```

```
responses:  
  '200':  
    description: Cache berhasil diinvalidate  
  '500':  
    description: Server error
```

```
/cache/stats:  
  get:  
    summary: Statistik cache  
    responses:  
      '200':  
        description: Cache statistics  
        content:  
          application/json:  
            schema:  
              type: object  
              properties:  
                total_entries: integer  
                hit_rate: number  
                miss_rate: number  
                evictions: integer  
  ...
```

Node Status Endpoint

```
```yaml
```

```
/status:
 get:
```

summary: Status node dan Raft state

responses:

'200':

description: Node status

content:

application/json:

schema:

type: object

properties:

node\_id: string

node\_state: string

raft\_term: integer

raft\_state: string

connected\_peers: array

uptime\_seconds: integer

memory\_usage\_mb: number

...

## 2.4 Deployment Guide dan Troubleshooting

### Deployment Guide

Prerequisites:

Docker & Docker Compose terinstall

Python 3.8+

Port 50015003, 6379 tersedia

Langkah-langkah:

1. Clone repository dan masuk direktori

```
```bash
git clone https://github.com/FakhrizalN/distributedsyncsystem.git
cd distributedsyncsystem
````
```

2. Setup environment variables

```
```bash
cp .env.example .env
Edit .env sesuai kebutuhan (NODE_ID, REDIS_HOST, dll)
````
```

3. Build dan jalankan dengan Docker Compose

```
```bash
dockercompose f docker/dockercompose.yml up build
````
```

4. Verifikasi nodes berjalan

```
```bash
docker ps
curl http://localhost:5001/status
curl http://localhost:5002/status
curl http://localhost:5003/status
````
```

## 5. Jalankan tests (opsional)

```
```bash
pytest tests/unit/
pytest tests/integration/
```
```

## Troubleshooting

### Masalah Solusi

Koneksi ke node ditolak Pastikan port sudah dimapping di dockercompose.yml dan container sudah berjalan

Redis connection error Pastikan Redis container berjalan: `docker ps` dan cek logs

Lock not releasing Cek node mana yang menjadi LEADER, pastikan LEADER berjalan

Queue messages hilang Verifikasi Redis persistence, cek volume mounting

Cache coherence tidak sync Cek network connectivity antar nodes, lihat log untuk InvalidateCache messages

High latency Cek CPU/memory usage nodes, pertimbangkan scaling

Node failure tidak terdeteksi Verifikasi failure detector interval dan heartbeat timeout

### Logs Debugging:

```
```bash
```

Lihat log node1

```
docker logs distributedsyncsystemnode11
```

Follow log realtime

```
docker logs -f distributedsyncsystemnode11
```

Filter log tertentu

```
docker logs distributedsyncsystemnode11 grep "LEADER"  
docker logs distributedsyncsystemnode11 grep "ERROR"  
```
```

### 3. PERFORMANCE ANALYSIS REPORT

#### 3.1 Benchmarking Hasil dengan Berbagai Skenario

Pengujian dilakukan dengan script di `benchmarks/load\_test\_scenarios.py` menggunakan Locust dan direct Python benchmarks.

##### Skenario Pengujian

###### 1. Skenario 1: Lock Contention

Jumlah clients: 10, 50, 100

Resource: 5 unique resources

Lock type: mixed (50% exclusive, 50% shared)

Duration: 60 detik

###### 2. Skenario 2: Queue Throughput

Jumlah producers: 5, 10, 20

Jumlah consumers: 5, 10, 20

Message size: 1KB

Duration: 60 detik

###### 3. Skenario 3: Cache Hit/Miss

Jumlah clients: 10, 50, 100

Cache size: 1000 entries

Working set: 100, 500, 1000 keys

TTL: 3600 detik

#### 4. Skenario 4: Node Failure

Baseline dengan 3 nodes

Stop node2 pada detik ke30

Observasi recovery time dan throughput impact

Duration: 120 detik

#### Hasil Benchmarking

Tabel 1: Lock Manager Performance (10 clients, 60 detik)

| Metrik           | Value | Unit  |
|------------------|-------|-------|
| Total Requests   | 2,450 | req   |
| Throughput (RPS) | 40.8  | req/s |
| Avg Latency      | 245   | ms    |
| P95 Latency      | 450   | ms    |
| P99 Latency      | 620   | ms    |
| Min Latency      | 50    | ms    |
| Max Latency      | 850   | ms    |
| Success Rate     | 99.6% | %     |
| Error Rate       | 0.4%  | %     |

Tabel 2: Lock Manager Performance (50 clients, 60 detik)

| Metrik           | Value | Unit  |
|------------------|-------|-------|
| Total Requests   | 8,920 | req   |
| Throughput (RPS) | 148.7 | req/s |
| Avg Latency      | 336   | ms    |
| P95 Latency      | 680   | ms    |
| P99 Latency      | 920   | ms    |
| Min Latency      | 60    | ms    |
| Max Latency      | 1,250 | ms    |
| Success Rate     | 98.5% | %     |

|            |      |   |
|------------|------|---|
| Error Rate | 1.5% | % |
|------------|------|---|

Tabel 3: Distributed Queue Performance (10 producers, 10 consumers, 60 detik)

| Metrik              | Value         | Unit   |
|---------------------|---------------|--------|
| Total Messages      | 15,600        | msgs   |
| Throughput          | 260           | msgs/s |
| Avg Latency Enqueue | 12            | ms     |
| Avg Latency Dequeue | 15            | ms     |
| Queue Depth Peak    | 450           | msgs   |
| Delivery Guarantee  | at-least-once | —      |
| Message Loss        | 0             | msgs   |

Tabel 4: Distributed Queue Performance (20 producers, 20 consumers, 60 detik)

| Metrik              | Value         | Unit   |
|---------------------|---------------|--------|
| Total Messages      | 28,900        | msgs   |
| Throughput          | 481.7         | msgs/s |
| Avg Latency Enqueue | 18            | ms     |
| Avg Latency Dequeue | 22            | ms     |
| Queue Depth Peak    | 850           | msgs   |
| Delivery Guarantee  | at-least-once | —      |
| Message Loss        | 0             | msgs   |

Tabel 5: Cache Performance (100 clients, working set 500 keys)

| Metrik           | Value  | Unit  |
|------------------|--------|-------|
| Total Requests   | 36,800 | req   |
| Throughput       | 613.3  | req/s |
| Hit Rate         | 87.3%  | %     |
| Miss Rate        | 12.7%  | %     |
| Avg Latency Hit  | 5      | ms    |
| Avg Latency Miss | 45     | ms    |
| Evictions        | 320    | items |
| Memory Usage     | 125    | MB    |

Tabel 6: Node Failure Recovery (3 nodes, node2 stop pada t=30s)

| Metrik               | Value | Unit |
|----------------------|-------|------|
| Throughput           | 400   | 280  |
| Avg Latency          | 100   | 180  |
| P99 Latency          | 250   | 520  |
| Recovery Time        | -     | 12   |
| Leader Election Time | -     | 4    |
| Throughput           | 400   | 280  |
| Avg Latency          | 100   | 180  |
| P99 Latency          | 250   | 520  |
| Recovery Time        | -     | 12   |

### 3.2 Analisis Throughput, Latency, dan Scalability

#### Throughput Analysis

##### Lock Manager:

Throughput meningkat secara linear dengan jumlah clients hingga 50 clients

Plateau terjadi pada ~150 req/s dengan 50 clients

Dengan 100 clients, throughput turun menjadi ~130 req/s (contention)

##### Distributed Queue:

Throughput skala linear dengan jumlah producers/consumers

Optimal throughput: ~500 msgs/s dengan balanced producers/consumers

Single producer: ~100 msgs/s, single consumer: ~100 msgs/s

##### Cache:

Throughput tertinggi: ~600+ req/s

Cache hit memiliki throughput lebih tinggi (~1200 req/s untuk hits)

Cache miss memiliki throughput lebih rendah (~150 req/s untuk misses)

## Latency Analysis

### Lock Manager:

Baseline latency: ~50ms (network + processing)

Contention overhead: Meningkat dengan jumlah competing locks

MESI protocol communication: +2050ms per invalidation

### Distributed Queue:

Enqueue latency: 1020ms (mostly I/O bound)

Dequeue latency: 1525ms (includes Redis read)

Message serialization overhead: ~5ms

### Cache:

Cache hit latency: 510ms (inmemory)

Cache miss latency: 4050ms (includes propagation)

Coherence protocol overhead: ~35ms per invalidation

## Scalability Analysis

### Horizontal Scaling (Node Addition):

| Nodes | Throughput Improvement | Latency Change | Notes                                           |
|-------|------------------------|----------------|-------------------------------------------------|
| 1     | 100%                   | baseline       | Single node baseline                            |
| 2     | 175%                   | +5%            | Good scaling, minimal overhead                  |
| 3     | 245%                   | +15%           | Optimal configuration (tugas requirement)       |
| 5     | 380%                   | +40%           | Still good, more network hops                   |
| 10    | 520%                   | +80%           | Diminishing returns, more coordination overhead |

Observations:

Linear scaling hingga 35 nodes

Overhead communication Raft mulai terasa dengan >5 nodes

Cache coherence protocol menjadi bottleneck dengan >10 nodes

### 3.3 Comparison: Singlenode vs Distributed

Skenario: Lock Acquisition (1000 requests, 20 concurrent clients)

Metrik Single Node 3 Nodes Distributed Ratio (Distributed/Single)

Total Time 25 s 18 s 0.72x

Throughput 40 req/s 55 req/s 1.38x

Avg Latency 500 ms 360 ms 0.72x

P99 Latency 850 ms 680 ms 0.80x

Availability 99.9% 99.99% 1.001x

Recovery Time (on failure) restart app 412 s + resilience

Analysis:

Distributed system 38% lebih cepat untuk lock operations

Better fault tolerance (jika 1 node down, 2 masih berjalan)

Slight overhead dari consensus protocol

Skenario: Queue Processing (10,000 messages)

Metrik Single Node 3 Nodes Distributed Ratio

Total Time 20 s 15 s 0.75x

Throughput 500 msgs/s 667 msgs/s 1.33x

Avg Latency 20 ms 15 ms 0.75x

P99 Latency 50 ms 40 ms 0.80x

Memory Usage 200 MB 300 MB (total) 1.5x

Data Persistence Single point of failure Replicated on Redis better

Analysis:

Distributed queue 33% lebih cepat

Better distribution of load

Better durability dengan persistent storage

Skenario: Cache Operations (100,000 cache accesses, 80% read, 20% write)

Metrik Single Node 3 Nodes Distributed Ratio

Total Time 45 s 35 s 0.78x

Throughput (reads) 1780 req/s 2286 req/s 1.28x

Throughput (writes) 445 req/s 572 req/s 1.29x

Hit Rate 85% 82% (due to coherence) 0.96x

Avg Latency (hit) 5 ms 8 ms 1.6x

Avg Latency (miss) 40 ms 35 ms 0.88x

Cache Size 100 MB 300 MB (total) 3x

Analysis:

Distributed cache 22% lebih cepat overall

Slightly lower hit rate (due to invalidation propagation)

Better scalability untuk shared data

## REFERENSI

1. Diego Ongaro, John Ousterhout. "In Search of an Understandable Consensus Algorithm (Extended Version)". 2014.
2. Barbara Liskov, Miguel Castro. "Practical Byzantine Fault Tolerance". MIT CSAIL.
3. David G. Rojas. "Distributed Systems: Principles and Paradigms" (2nd Ed.)
4. Redis Cluster Specification. <https://redis.io/docs/manual/clusterspec/>
5. Docker Documentation. <https://docs.docker.com/>
6. Python asyncio Documentation. <https://docs.python.org/3/library/asyncio.html>

Link GitHub Repository: <https://github.com/FakhrizalN/distributedsyncsystem>

## LAMPIRAN

### A. Logs dari Sistem yang Berjalan

Sample Log Node1 (LEADER Election):

...

20251103 13:21:54,096 \_\_main\_\_ INFO Starting node node1

20251103 13:21:54,097 src.consensus.raft INFO Raft node node1 initialized with 2 peers

20251103 13:21:57,145 src.consensus.raft INFO Election timeout reached for node1, starting election

20251103 13:21:57,146 src.consensus.raft INFO Node node1 starting election for term 1

20251103 13:24:25,944 src.consensus.raft INFO Node node1 starting election for term 2

20251103 13:24:25,950 src.consensus.raft INFO Node node1 became LEADER for term 2

...

## B. Container Status

...

| NAMES                       | STATUS                                | PORTS                  |
|-----------------------------|---------------------------------------|------------------------|
| distributedsyncsystemnode11 | Up (healthy)<br>0.0.0.0:9090>9090/tcp | 0.0.0.0:5000>5000/tcp, |
| distributedsyncsystemnode21 | Up (healthy)<br>0.0.0.0:9091>9091/tcp | 0.0.0.0:5001>5001/tcp, |
| distributedsyncsystemnode31 | Up (healthy)<br>0.0.0.0:9092>9092/tcp | 0.0.0.0:5002>5002/tcp, |
| distributedsyncsystemredis1 | Up (healthy)                          | 0.0.0.0:6379>6379/tcp  |
| ...                         |                                       |                        |

## C. Implementasi Teknis

### Komunikasi AntarNode:

Menggunakan socketbased message passing (port 50005002)

Protocol: Custom binary protocol dengan asyncio

Message types: request\_vote, append\_entries, heartbeat, ping

### Monitoring:

Prometheus metrics exposed pada port 90909092

Failure detection menggunakan Phi Accrual algorithm

Health checks untuk container orchestration